

BRAJMOHAN RAJPUT

AI/ML Engineer

+91-7489577065 brajrajput4@gmail.com LinkedIn GitHub Portfolio

PROFESSIONAL SUMMARY

AI/ML Engineer (B.Tech, 2027) with hands-on experience building production-ready Generative AI applications, cloud-native systems, and end-to-end ML pipelines. Skilled in conversational RAG, LLM-powered FastAPI backends, Kubernetes deployments with ELK Stack observability, and scalable MLOps practices using Docker, AWS, and GitHub Actions.

TECHNICAL SKILLS

Generative AI & LLMs: GPT-4, LLM Applications, Retrieval-Augmented Generation (RAG), LangChain, Prompt Engineering, HuggingFace Embeddings, ChromaDB, Conversation Memory

Machine Learning: Scikit-learn, Feature Engineering, Model Training & Evaluation, Hyperparameter Tuning (GridSearchCV), Classification, Regression

Backend Development: Python, FastAPI, REST APIs, Streamlit

Cloud & DevOps: AWS (S3, EC2, Lambda, CodePipeline), Azure AI Foundry, Docker, Kubernetes (Minikube), GitHub Actions, CI/CD Pipelines

Logging & Observability: ELK Stack (Filebeat, Logstash, Elasticsearch, Kibana), Centralized Log Management, Real-Time Monitoring Dashboards, MLflow

Data Engineering: Apache Airflow, Apache Spark, PostgreSQL, SQL, ETL/ELT Pipelines

Developer Tools: Git, GitHub, Linux, VS Code

PROJECTS

Source Code Analyser (Production-Ready Conversational RAG)

Apr 2026

Python, FastAPI, Streamlit, LangChain, ChromaDB, HuggingFace Embeddings, Groq API

- Built a production-ready conversational RAG system that indexes entire GitHub repositories into a ChromaDB vector store, enabling semantic code search and repository-aware Q&A over complex codebases.
- Architected a modular FastAPI backend with LangChain, isolating retrieval, vector storage, session management, and rolling conversation summarization into dedicated, independently testable components.
- Reduced prompt token usage in extended sessions by implementing session-scoped memory with automatic conversation summarization, preserving long-term context without exceeding LLM context limits.

AI Powered Job Analyzer — ELK Stack & Kubernetes

Jun 2026

Python, GPT-4 API, FastAPI, Streamlit, Filebeat, Logstash, Elasticsearch, Kibana, Docker, Kubernetes (Minikube)

- Engineered a cloud-native GenAI application that leverages GPT-4 to automatically screen resumes against job descriptions, producing match percentages, skill gap analysis, and hiring decision summaries.
- Deployed the application as a containerized pod on a Minikube Kubernetes cluster inside a VM, demonstrating production-grade container orchestration with horizontal scaling capability.
- Implemented a full ELK Stack observability pipeline using Filebeat as a sidecar agent to ship application logs to Logstash for processing, Elasticsearch for indexed storage, and Kibana for real-time system health dashboards.

Two-Stage Loan Approval & Valuation ML System

Nov 2025

Python, Scikit-learn, Streamlit Cloud, GridSearchCV

- Architected a two-stage ML inference pipeline that chains a classification model (loan approval) with a regression model (loan amount estimation), enabling a unified, sequential prediction workflow in a single user request.
- Engineered end-to-end Scikit-learn Pipelines with modular preprocessing and feature engineering stages; applied GridSearchCV hyperparameter tuning to optimize model performance across both stages.
- Deployed a production-style interactive ML application on Streamlit Cloud, demonstrating full-cycle capability from data preprocessing and model training to live inference and user-facing UI.

CERTIFICATIONS

Certified LLM Security Professional (CLLMSP) | Red Team Leaders

June 2026

Credential ID: 93cc00f8370b7800

EDUCATION

Bachelor of Technology (B.Tech)

S.G.S.I.T.S., Indore

Expected: 2027

CGPA: 7.0 / 10